

Explaining Before Installing: An Explainable AI Privacy Risk Report that Improves User Comprehension of Android Application Behaviour

Rafid Ahmed

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
rafidahmed816@gmail.com

Muntasir Hossain

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
muntasirhossain2003@gmail.com

Ifti Akkas

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
iftiakkas@iut-dhaka.edu

Rubaeyat Hossain

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
rubaeyat@iut-dhaka.edu

Zawad Bin Shoukat

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
zawadtasinmn@gmail.com

Syed Huzzatullah Mihad

Department of Computer Science and
Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
huzzatullah@iut-dhaka.edu

Hasan Mahmud

Corresponding Author
SSL, Department of Computer Science
and Engineering
Islamic University of Technology
(IUT), OIC
Dhaka, Bangladesh
hasan@iut-dhaka.edu

ABSTRACT

Application stores disclose what data an application collects, but they do not tell a user whether that collection is reasonable. The Android *Data safety* section, for example, lists categories such as “Personal info” and “Device or other IDs” without conveying how much privacy risk the combination represents or why. Users are therefore given transparency without interpretability. We present an explainable AI (XAI) based privacy risk assessment and decision support system that converts static application features into a graded privacy risk score, uses SHAP to identify the features responsible for that score, and renders those factors as a short, plain-language justification shown before installation. We evaluated the resulting interface against the current Google Play *Data safety* page in a within-subjects study with 34 participants. Comprehension of the privacy risk rose from $M = 2.15$ to $M = 3.79$ on a five-point scale (Wilcoxon signed-rank, $p < .001$, $r_{rb} = 0.80$), and the ability to understand the information without technical knowledge rose from $M = 2.18$ to $M = 3.85$ ($p < .001$, $r_{rb} = 0.82$). In forced choice, 31 of 34 participants preferred the explanatory interface on every comparative item ($p < .001$). The contribution of this work is not a new malware detector but the demonstration that an explanation layer placed between a risk model and an end user measurably changes what non-expert users understand about

an application before they install it, together with design guidance on how brief such explanations must be to remain useful.

KEYWORDS

Human-Centered AI, Interaction Design, Usability, Explainable AI, HCI, Android Privacy, Decision Support

1 INTRODUCTION

Mobile applications routinely request access to contacts, precise location, messages, device identifiers and network services. Both major application stores now require developers to declare these practices, and Android surfaces them in a *Data safety* section on every store listing. The disclosure problem is, in a narrow sense, solved: the information is present, standardised and easy to find.

The interpretation problem is not. A user reading that an application collects “Personal info,” “Device or other IDs,” “App activity” and “Location” receives four category labels and no basis on which to judge them. The labels do not say whether this combination is ordinary for the type of application being installed, which of the declared behaviours matters most, or what the practical consequence of the combination might be. Felt et al. [8] showed more than a decade ago that only a small minority of users both notice and correctly interpret Android permission information, and later work on privacy labels reports that structured disclosures still fail

to answer the questions users actually have [28]. Disclosure has improved; comprehension has not followed.

Two lines of work address parts of this gap without closing it. Privacy nutrition labels [11] and personalised privacy assistants [16] make privacy information more legible and more actionable, but they still present *what* is collected rather than *how concerning* it is. Machine learning work on Android security [3, 25, 27] produces accurate risk and malware classifiers, but their output is a label or a probability aimed at analysts, and a user told only that an application scores 0.84 is no better placed to decide than one reading a category list. Neither tradition supplies the missing piece: a justification, in ordinary language, of why this particular application is being flagged.

We therefore treat the problem as one of human-centred AI rather than of classification accuracy. Our system assesses privacy risk from static application features, attributes that assessment to specific features using SHAP [18], and translates the attributed features into a two-to-three sentence justification presented alongside a graded risk score at the moment of the install decision. The risk judgement remains with the model; the language layer is constrained to restate what the model and the attribution actually found, never to reach an independent verdict. The user remains the decision maker, and the interface is explicit that an automated assessment indicates potential concern rather than proven misuse.

This paper makes three contributions.

- (1) **An explanation architecture for privacy risk.** We describe a pipeline—feature extraction, gradient-boosted risk assessment, SHAP attribution, constrained natural-language generation, decision-support interface—in which every user-visible sentence is traceable to a model attribution rather than generated freely.
- (2) **Empirical evidence that the explanation layer changes comprehension.** In a within-subjects study with 34 participants, the explanatory interface produced large, statistically significant improvements in both understanding of the risk and understanding without technical background, relative to the deployed Google Play *Data safety* page.
- (3) **Design implications for explanation length and form.** Participant feedback converged on a constraint that the quantitative results alone would have hidden: the explanation must be short and scannable. Prose that is more complete but longer was explicitly rejected (“I don’t want to read a novel”), which sharpens existing XAI guidance for consumer-facing, pre-decision contexts.

2 RELATED WORK

2.1 Communicating Mobile Privacy to Non-Experts

Research on Android permissions has consistently found a comprehension rather than an exposure problem. Felt et al. [8] reported that most users neither attended to nor understood install-time permission requests. Lin et al. [14] showed that user concern depends less on which permission is requested than on whether the request matches the expected purpose of the application—a location request from a maps application and the same request from a flashlight

application are judged very differently. This finding motivates a central design decision in our work: the explanation is written in terms of what the combination of behaviours enables, not as an enumeration of permission strings.

Structured disclosure formats improve legibility. The privacy nutrition label [11] demonstrated that a standardised layout speeds accurate information retrieval, and Kelley et al. [12] showed that privacy information placed in the application selection flow shifts choices. Nudging studies report similar effects in the field [1]. However, evaluations of the deployed Data safety labels find that they still do not answer users’ actual questions [28]. All of these formats share a structural property: they report *what* is collected. Our system adds an assessment of *how concerning* the collection is, and a reason for that assessment.

2.2 Explainable AI and Its Human Evaluation

Feature-attribution methods such as LIME [21] and SHAP [17, 18] identify the inputs responsible for a prediction, and are the standard choice for tree ensembles. Their limitation for our purpose is that their native output—a ranked list of feature names with signed contributions—is a technical artefact. Lipton [15] and Doshi-Velez and Kim [7] argue that interpretability claims must be evaluated against a stated purpose and audience rather than asserted, and Miller [19] draws on social science to show that human explanation is selective and contrastive rather than exhaustive. Kaur et al. [9] found that even data scientists over-trust attribution tools they do not fully understand, and Bansal et al. [4] report that explanations can increase acceptance of AI output without improving decision quality.

These results shape both our design and our claims. We do not present raw SHAP values to users; we use them as the content constraint for a short justification. We also deliberately restrict our claim to comprehension and preference, which our study measured, rather than to decision correctness, which it did not.

2.3 Human-Centered AI Principles

Shneiderman’s framework [23, 24] argues for systems that are simultaneously highly automated and highly controllable, and the Guidelines for Human-AI Interaction [2] specify concrete obligations: make clear what the system can do, make clear how well it does it, and support efficient correction. Liao et al. [13] derive an XAI question bank from practitioner interviews, in which “why this outcome” and “what would change it” dominate. Interaction design texts [22] and goal-directed design [6] supply the complementary requirement that the artefact be organised around a user goal—here, deciding whether to install—rather than around the structure of the underlying computation.

2.4 Research Gap

Existing work either makes privacy disclosures more legible without assessing them, or assesses application risk accurately without explaining the assessment to a non-expert. The gap we address is the junction: an assessment that is attributed, translated into ordinary language, bounded in its claims, and placed at the point of decision—and, critically, an empirical test of whether that junction changes what ordinary users understand.

3 METHOD (INTEGRATING ASSIGNMENT 1 AND 2)

This section integrates the requirements analysis of Assignment 1 with the design and prototyping work of Assignment 2, and shows how the evidence gathered about users determined the form of the final interface.

3.1 Requirement Analysis (Assignment 1)

3.1.1 Users and Stakeholders. The primary users are ordinary Android users who install applications from an application store and have no security or machine-learning background. Secondary stakeholders include application store operators, who bear the consequence of misleading risk communication, and application developers, whose products may be penalised by an unclear or unjustified warning. The distinction matters for design: a system that labels applications without justification transfers reputational risk to developers without giving users a basis to evaluate the label.

3.1.2 Data Gathering. We used a structured questionnaire administered online, supplemented by inspection of the deployed Google Play *Data safety* page as an artefact analysis. A questionnaire was chosen over interviews because the research question concerns comparative comprehension across a population of non-expert users rather than the depth of any individual’s mental model, and because a between-interface comparison requires a consistent stimulus for every respondent. The artefact analysis established the baseline condition and fixed the vocabulary—“Data shared & collected”, “Personal info”, “Device or other IDs”—that the proposed interface would have to improve upon.

3.1.3 Analysis Method. Closed items were analysed descriptively and comparatively. Free-text responses were analysed with thematic analysis following Braun and Clarke [5]: responses were coded independently, codes were grouped into candidate themes, and themes were reviewed against the full response set. With a small free-text corpus this yields orientation rather than saturation, and we treat the themes accordingly.

3.1.4 Personas and Scenarios. Persona 1 — Rifat, 22, undergraduate student. Installs applications several times a month, mostly on recommendation. Scrolls past the Data safety section because the categories mean nothing to him concretely. Goal: avoid installing something that will quietly harvest his contacts, without having to become an expert. Frustration: every application appears to collect everything, so the disclosure gives him no way to discriminate.

Persona 2 — Nusrat, 27, working professional. Privacy-conscious and moderately technical. Reads the permission list and can name what each permission does, but cannot judge whether a given *combination* is excessive for the application’s stated purpose. Goal: a defensible reason to accept or reject an application in under a minute.

Scenario. Rifat finds a note-taking application and a messaging application. Both declare access to device information and messaging functions. The store presents both disclosures in the same neutral format. He needs to know that in one case the combination is routine and in the other it is not, and he needs that distinction before, not after, installation.

Table 1: Requirements–goal traceability matrix.

Req	User goal	Interface element
FR1	Judge severity at a glance	Risk badge, score /100, gauge bar
FR2	Know what drove the judgement	Content of “Why this assessment?”
FR3	Understand without expertise	Plain-language justification text
FR4	Avoid over-trusting the system	“Important note” disclaimer block
FR5	Verify independently	“Review app information” link
FR6	Decide before committing	Placement on the pre-install screen
NFR1	Understand without expertise	No permission strings in body text
NFR2	Decide quickly	Single screen, no scrolling to reach verdict
NFR3	Calibrate concern	Green/amber/red used only on the risk band
NFR4	Trust the reason given	Generation constrained to SHAP output

3.1.5 Task Analysis. The pre-install privacy decision decomposes into four steps: (1) locate privacy information on the store listing; (2) interpret the declared behaviours; (3) judge whether those behaviours are proportionate to the application’s purpose; (4) decide to install or seek an alternative. The deployed interface supports step 1 well and steps 2–3 poorly. Our design targets steps 2 and 3 specifically, and leaves step 4 to the user.

3.1.6 Requirements. Functional requirements.

- FR1 Produce a graded privacy risk level (Low / Medium / High) with a bounded numeric score for an application.
- FR2 Identify the application features that contributed most to that assessment.
- FR3 Present those contributing factors as a short, non-technical justification.
- FR4 State explicitly that the assessment is automated and indicative, not a determination of malicious behaviour.
- FR5 Provide access to the underlying detailed application information for users who want it.
- FR6 Present the assessment before the install decision is taken.

Non-functional requirements.

- NFR1 The justification must be comprehensible without security or machine-learning knowledge.
- NFR2 The assessment must be readable within a short glance, consistent with store browsing behaviour.
- NFR3 Visual risk encoding must be consistent and non-alarmist.
- NFR4 Every user-visible claim must be traceable to a model attribution.

3.1.7 Requirements–Goal Traceability. Table 1 maps each requirement to the user goal that motivates it and to the interface element that satisfies it.




Figure 1: Interface A (baseline): the deployed Google Play *Data safety* page. Collection categories are disclosed; no assessment, ranking or justification is provided.

3.2 Design and Prototyping (Assignment 2)

3.2.1 Design Rationale. Four decisions follow directly from the requirements above.

A graded score rather than a binary verdict. Step 3 of the task analysis is a proportionality judgement, which a binary safe/unsafe label cannot support. A 0–100 score with three named bands (FR1) allows a user to distinguish a routine application from a borderline one.

Explanation in terms of capability, not permission. Lin et al. [14] show that users reason about purpose, not permission identifiers. The justification therefore refers to “messaging functions” and “device identification” rather than SEND_SMS or READ_PHONE_STATE, satisfying NFR1.

An explicit epistemic boundary. Given evidence that explanations can inflate trust beyond warrant [4], every screen carries a disclaimer that an automated assessment indicates potential risk and does not prove misuse (FR4), in line with the Human-AI Interaction guideline to make clear how well the system performs [2].

Verdict above the fold, detail on demand. NFR2 requires that the risk level be readable without scrolling; FR5 requires that the underlying detail remain reachable. The layout therefore places badge, score and gauge at the top and defers full application information to a link.

3.2.2 Interaction Flow. The user opens an application listing, the assessment is computed and displayed, the user reads the risk band and, if they wish, the justification, optionally expands the detailed application information, and then installs or abandons. The interface does not block installation at any point; it informs it.

3.2.3 Prototype. Figure 1 shows the baseline condition—the deployed Google Play *Data safety* page. Figure 2 shows the proposed interface across three risk bands.

3.3 Human-Centered AI Considerations

Explainability. The user is never shown a bare score. The “Why this assessment?” block names the behavioural signals that produced the score and states what their combination could enable,

addressing the “why this outcome” question that dominates practitioner XAI needs [13].

Human control. The system is advisory. It does not block, rank or remove applications, and the install decision is never automated. The “Review app information” affordance lets a user inspect the evidence and disagree with the assessment.

Trust calibration. Rather than maximising trust, the design bounds it. The disclaimer is persistent and appears even on the Low-risk screen, so that the limits of the assessment are communicated in the case where the user has least incentive to read them.

4 SYSTEM DESIGN AND IMPLEMENTATION

4.1 System Overview

The system is a five-stage pipeline. An Android application package supplies static features; a gradient-boosted classifier produces a privacy risk assessment; SHAP attributes that assessment to individual features; a constrained language layer converts the top-ranked attributions into a short justification; and the decision-support interface presents score, justification and disclaimer together before installation.

The separation between the training corpus and the runtime input is deliberate. The model is trained on a labelled corpus of Android application feature vectors, whereas at run time features are extracted from an actual application package into the identical feature space. The user is never shown the training data; they are shown an assessment of the application in front of them.

4.2 Functional Design and Features

Graded risk assessment (FR1). The classifier output is mapped to a 0–100 score and to one of three bands. The gauge bar makes the position within the band visible, so that a score of 51 does not read identically to a score of 79.

Factor identification (FR2). SHAP values rank the features by contribution to the individual prediction. Only the highest-contributing features enter the explanation, consistent with the finding that human explanation is selective rather than exhaustive [19].

Plain-language justification (FR3, NFR1, NFR4). The language layer receives the risk band and the ranked attributions and is constrained to restate them. It does not have access to an independent judgement of the application and cannot introduce a factor the model did not identify. This constraint is what makes NFR4 checkable: every clause in the justification corresponds to an attributed feature.

Epistemic disclaimer (FR4). Fixed text, not generated, so that it cannot be weakened by the language layer.

Detail on demand (FR5). “Review app information” exposes the underlying declared behaviours for users who want the primary evidence.

4.3 Interface Design and Interaction Flow

The screen is organised top to bottom in decreasing order of decision value: application identity, scope statement, risk band and score, justification, disclaimer, detail link. A user who reads only the first third of the screen still obtains the verdict; a user who reads all of it obtains the reason and its limits. Colour is applied only to the




Figure 2: Interface B (proposed): the explainable privacy risk report at Low (22/100), Medium (51/100) and High (84/100) risk. Each screen pairs a graded score with a “Why this assessment?” justification derived from the model’s feature attributions, an “Important note” stating the limits of an automated assessment, and a link to the detailed application information.

risk band and gauge, so that the encoding remains unambiguous and the rest of the screen stays neutral (NFR3).

4.4 Human-Centered AI Implementation

AI performs assessment and attribution only. It does not act, filter or decide. Communication of the assessment is split across three registers—a numeric score for magnitude, a colour band for immediate severity, and prose for reasoning—so that a user who does not read the prose still receives a calibrated signal. Override is implicit but total: the user installs or does not.

4.5 Implementation Details

The risk model is a LightGBM gradient-boosting classifier [10], selected because the application representation is high-dimensional, sparse and binary—conditions under which gradient-boosted trees are strong and, equally important, under which exact tree-based SHAP values are tractable [17]. Attribution uses the shap library’s TreeExplainer; data handling and evaluation use scikit-learn [20].

The model is trained on a public corpus of static Android application features, in which each application is represented as a binary vector over requested permissions, declared intents and API-related behaviours.

Implementation limitations. The analysis is static. Behaviour that appears only at run time, code loaded dynamically, and obfuscated or packed payloads are outside its reach, and an application can therefore present a benign static profile while behaving otherwise. The system is positioned accordingly as a privacy risk indicator, not a malware detector—a boundary the interface states to the user rather than leaving implicit.

5 EVALUATION

5.1 Experimental Design

Participants. Thirty-four participants completed the study. Table 2 reports the sample. The majority were aged 21–25 (79.4%), with 11.8% aged 18–20 and 8.8% aged 26–30. Self-rated familiarity with Android permissions and privacy information averaged $M = 3.29$

Table 2: Participant demographics and self-reported background ($N = 34$).

Variable	Level	n (%)
Age group	18–20	4 (11.8)
	21–25	27 (79.4)
	26–30	3 (8.8)
Install frequency	Rarely	5 (14.7)
	Occasionally	13 (38.2)
	Frequently	11 (32.4)
	Very frequently	5 (14.7)
Permission familiarity	1 (none)	1 (2.9)
	2	6 (17.6)
	3	14 (41.2)
	4	8 (23.5)
	5 (high)	5 (14.7)

($SD = 1.02$) on a five-point scale, spanning the full range from 1 to 5, so the sample contains both low- and high-familiarity respondents rather than only novices. Participants were recruited by convenience sampling and were not compensated.

Design. The study used a within-subjects design with interface as the independent variable at two levels: Interface A, the deployed Google Play *Data safety* page (Figure 1), and Interface B, the proposed explainable privacy report (Figure 2). A within-subjects design was chosen because the dependent measures are subjective comprehension ratings, which vary substantially between individuals; having each participant rate both interfaces removes that between-person variance and permits a direct paired comparison, which is valuable at this sample size.

Dependent variables. Two five-point Likert items were collected for each interface: ease of understanding the privacy risk presented, and ease of understanding the information without technical knowledge. Three forced-choice items then asked which interface made the risk easier to understand, which explained the reason behind the risk more clearly, and which the participant would prefer in practice. A free-text item asked for the main improvement to the proposed interface.

Context. The evaluation was conducted remotely through an online questionnaire, with both interfaces presented as static screen images. This is a comprehension study of the presentation, not a usability study of a live system; no interaction was logged.

5.2 Procedure

Participants first answered demographic and background questions. They were then shown Interface A, described as the privacy information currently shown on an application store listing, and rated it on the two comprehension items. They were then shown Interface B across its three risk bands, described as a proposed alternative, and rated it on the same two items. The three forced-choice comparisons followed, with both interfaces available for reference, and the questionnaire closed with the free-text improvement item. No training was given, which is deliberate: the interface is intended for users encountering it cold on a store listing, and any briefing would have measured the briefing rather than the design.

Table 3: Paired comparison of comprehension measures ($N = 34$, five-point scale). Wilcoxon signed-rank test; r_{rb} is the matched-pairs rank-biserial correlation.

Measure	A: M (SD)	B: M (SD)	Test
Understanding of privacy risk	2.15 (1.17)	3.79 (1.21)	$p < .001$ $r_{rb} = 0.80$
Understanding without technical knowledge	2.18 (1.10)	3.85 (1.06)	$p < .001$ $r_{rb} = 0.82$

5.3 Evaluation Metrics

Because the study evaluates comprehension of a static presentation rather than performance on an interactive system, the conventional efficiency measures—task completion time, error rate, number of interactions—do not apply and were not collected. The measures used were:

Quantitative. (i) Perceived comprehension of the privacy risk; (ii) perceived comprehension without technical knowledge; (iii) proportion of participants selecting each interface on three forced-choice items. Measures (i) and (ii) map directly to NFR1 and to steps 2–3 of the task analysis, which is where the baseline was judged weakest. Measure (iii) captures revealed preference, which is a stricter test than a rating because it forces a comparison.

Qualitative. Free-text responses on the main improvement needed, analysed thematically [5].

5.4 Analysis Approach

Likert responses are ordinal and paired, so the two interfaces were compared with the Wilcoxon signed-rank test [26] rather than a paired t -test. Effect size is reported as the matched-pairs rank-biserial correlation r_{rb} , on which 0.5 is conventionally regarded as a large effect. Forced-choice outcomes were tested against chance with a two-sided exact binomial test on the participants expressing a preference, with ties excluded. The significance threshold was $\alpha = .05$. Free-text responses were coded and grouped into themes, which are reported with counts to keep their weight visible given the small corpus.

6 RESULTS

6.1 Quantitative Results

Table 3 reports the paired comparison on both comprehension measures. Interface B was rated substantially higher on each.

Understanding of the privacy risk rose from $M = 2.15$ ($SD = 1.17$) for Interface A to $M = 3.79$ ($SD = 1.21$) for Interface B, a difference of 1.64 scale points ($W = 40.0$, $p < .001$, $r_{rb} = 0.80$, 26 non-tied pairs). Understanding without technical knowledge rose from $M = 2.18$ ($SD = 1.10$) to $M = 3.85$ ($SD = 1.06$), a difference of 1.67 points ($W = 32.5$, $p < .001$, $r_{rb} = 0.82$, 26 non-tied pairs). Both effects are large by conventional interpretation.

The shape of the distributions is as informative as the means. For Interface A, 22 of 34 participants (64.7%) gave a rating of 1 or 2 for understanding the risk, and only 4 (11.8%) gave 4 or 5. For Interface B the pattern inverts: 22 participants (64.7%) gave 4 or 5,

Table 4: Forced-choice comparisons ($N = 34$). Two-sided exact binomial test on the 33 participants expressing a preference.

Item	B	A	Tie	p
Easier to understand the risk	31	2	1	< .001
Explained the reason more clearly	31	2	1	< .001
Would prefer in practice	31	2	1	< .001

and 7 (20.6%) gave 1 or 2. The baseline was not merely rated lower on average; it failed for a clear majority of the sample.

Table 4 reports the forced-choice items. On all three, 31 of 34 participants selected Interface B, 2 selected Interface A, and 1 recorded a tie or no preference. Among the 33 participants expressing a preference, this is significant against chance on each item (two-sided exact binomial, $p < .001$).

The three participants who did not prefer Interface B on the overall preference item are worth isolating rather than averaging away. One rated Interface A at the ceiling (5 on both measures) and Interface B at the floor, and supplied the free-text comment quoted below; this participant also reported the highest install frequency and high familiarity. The pattern is consistent with an experienced user for whom the baseline is already sufficient and additional prose is overhead.

6.2 Qualitative Findings

Seven participants supplied free-text responses. Thematic analysis produced three themes.

Theme 1: Explanation length is a cost (2 responses). The most pointed response was “I want to know what information they collect. I don’t want to read a novel.” A second asked that the interface be made “more convenient and easier or user friendly.” Both locate the problem in the volume of prose rather than its content.

Theme 2: Structure over prose (1 response). “The assessment why some app is risky could be written in bullet points to make it more clear visually.” This is a concrete alternative to the paragraph form: the same attributed factors, rendered as a scannable list.

Theme 3: Approval with a clarity reservation (4 responses). “It’s all good”; “UI is good. But should be more clear.”; “UI”; “N/A”. These endorse the direction while, in two cases, repeating that clarity can improve.

No participant questioned the accuracy of the risk scores or asked how they were computed—a silence we return to in the Discussion.

6.3 Comparative Observations

The two conditions differ in kind, not degree. Interface A discloses categories; Interface B discloses categories, assesses them, ranks the contributing factors and justifies the assessment. The comprehension gap of roughly 1.65 scale points, together with near-unanimous preference, indicates that the addition is doing work. The near-identical results on the two comprehension measures (1.64 and 1.67 points; $r_{rb} = 0.80$ and 0.82) suggest that for this sample, understanding the risk and understanding it without technical background were not separable—the baseline’s failure was a single failure of interpretability, not two distinct ones.

6.4 Summary of Key Findings

What worked. The graded score with a plain-language justification produced large, significant gains on both comprehension measures and was preferred by 91.2% of participants on every comparative item.

What did not. The paragraph form of the justification was the one consistently criticised element. Participants wanted the same information in less space and in scannable structure.

Unexpected. No participant asked how the risk score was derived, or expressed doubt about it, despite a persistent on-screen disclaimer stating the assessment is automated and indicative. The explanation appears to have satisfied curiosity about the verdict without prompting scrutiny of the mechanism.

7 DISCUSSION

7.1 Why the Explanatory Interface Won

The baseline’s weakness is not that it withholds information but that it stops one step short of the user’s actual task. Step 3 of our task analysis—judging proportionality—requires a comparison between what an application does and what its purpose warrants, and a category list supplies only one side of that comparison. Interface B supplies the other side: it states which behaviours drove the assessment and what their combination could enable. This is consistent with Lin et al. [14]: users reason about purpose fit, and an interface that speaks in those terms is doing less translation work for them.

The equality of the two effect sizes supports this reading. If the baseline’s problem had been technical vocabulary alone, we would expect the “without technical knowledge” measure to improve more than the general comprehension measure. They improved almost identically, which suggests the missing element was interpretation itself rather than terminology—a gap that better wording of a category list could not close.

7.2 Which Design Decisions Held

The graded score and colour band held: participants oriented to the verdict without complaint, and no free-text response asked for a different risk encoding. The disclaimer held in the weak sense that no participant reported confusion about the system’s authority, though as noted below its effect on trust was not measured.

The paragraph justification did not hold. It satisfied the requirement to explain (FR3) while violating the requirement to be readable at a glance (NFR2), and participants detected the tension immediately. Our design treated “explanation” and “brevity” as compatible; the data indicates they trade off, and that in a pre-install context brevity has the stronger claim. The bullet-point suggestion from Theme 2 is the natural resolution and preserves the attribution chain, since the underlying SHAP ranking is already a list—it was the rendering into prose that added the length.

7.3 The Silence About the Model

That no participant questioned how the score was computed is the finding we regard with most caution. A benign reading is that the explanation answered the question users had, making the mechanism irrelevant to them. A less comfortable reading is that a fluent

justification licensed acceptance of a number without scrutiny—precisely the pattern Bansal et al. [4] report, where explanation increases acceptance without improving decision quality, and Kaur et al. [9] observe even among experts. Our study cannot distinguish these, because it measured comprehension and preference but not calibration: we never showed participants an incorrect assessment to see whether they would catch it. Until that test is run, the comprehension gain should not be read as evidence of appropriate trust.

7.4 The Dissenting Participants

The participant who preferred Interface A rated their own familiarity at the ceiling and installs applications very frequently. For such a user, the category list is already interpretable and the justification is redundant text between them and the decision. This suggests that explanation depth should be responsive to user expertise rather than fixed—a progressive disclosure design in which the verdict is always visible and the justification expands on demand would serve both this participant and the majority, and would simultaneously address the length complaint.

8 IMPLICATIONS FOR DESIGN

1. For consumer pre-decision contexts, explanation length is a usability constraint, not a presentation detail. XAI guidance emphasises selecting and contrasting [19], but our participants applied a harder budget: two to three sentences of prose were already experienced as too much at the moment of an install decision. Designers should treat the explanation's word count as a requirement with a stated limit, tested like any other.

2. Render attributions as structure, not prose. Feature attribution produces a ranked list. Converting it into paragraphs adds reading time without adding information. Keeping the list form—short labelled bullets in the user's vocabulary—preserves scannability and shortens the distance between what the model computed and what the user reads.

3. Make explanation depth progressive. A fixed explanation depth cannot serve both a novice and an expert. Verdict always visible, justification one tap away, evidence two taps away, satisfies both without penalising either.

4. Comprehension gains are not trust calibration, and should not be reported as such. Our study shows users understood the assessment better; it does not show they would have rejected a wrong one. Any system in this space should include at least one condition with a deliberately incorrect assessment, or its evaluation cannot speak to over-reliance.

5. State the epistemic boundary in fixed, non-generated text. When a language layer produces the explanation, the limits of the system should not be within its reach to soften. Holding the disclaimer outside the generated region is a small architectural decision with a direct effect on what the system can be made to claim.

9 LIMITATIONS

Sample. Thirty-four participants, recruited by convenience, of whom 79.4% were aged 21–25 and the majority of student background. The sample is younger and more technically exposed than

the Android user population, and the results should not be generalised to older or less digitally fluent users, who may need more explanation rather than less.

Stimulus and measures. Participants rated static images, not a working system. All quantitative measures are self-reported perceived comprehension; we did not test objective comprehension—for example by asking participants to state what an application could do after reading each interface. Perceived and actual understanding can diverge, and a fluent explanation is exactly the condition under which they are most likely to. Because the baseline used a real, recognisable application (Facebook) while the proposed interface used fictional applications, brand-related priors may have affected the baseline ratings; matching the applications across conditions would remove this confound.

Order effects. All participants saw Interface A before Interface B. Presentation order was not counterbalanced, so part of the difference may reflect order or contrast rather than the design alone.

Instrumentation. Two of the four Likert items were mislabelled in the deployed questionnaire, both referring to Interface B where the first of each pair concerned Interface A. The intended mapping was confirmed and the response distributions are internally consistent with it, but the questionnaire text as administered was ambiguous and should be corrected before reuse.

Free-text corpus. Seven responses support orientation, not thematic saturation. The length theme is stated by two participants; its prominence in our design implications reflects its specificity and its convergence with the third theme's clarity reservations, not its frequency.

System scope. The analysis is static, and the risk model's predictive performance is reported elsewhere in this paper as a design property rather than validated against a held-out benchmark within this study. The user evaluation tested the presentation of an assessment, not the correctness of the assessment.

10 CONCLUSION

We built an explainable AI privacy risk assessment and decision support system for Android applications, which converts static application features into a graded risk score, attributes that score to specific behavioural features using SHAP, and presents the result as a short plain-language justification at the point of the install decision. Evaluated against the deployed Google Play *Data safety* page with 34 participants, the interface improved perceived comprehension of privacy risk by 1.64 points on a five-point scale ($p < .001$, $r_{rb} = 0.80$) and comprehension without technical knowledge by 1.67 points ($p < .001$, $r_{rb} = 0.82$), and was preferred by 31 of 34 participants on every comparative item.

The contribution is the demonstration that an explanation layer between a risk model and an end user changes what non-experts understand before they install, together with the design constraint our participants imposed on it: the explanation must be brief and scannable, or it becomes the new obstacle. Attribution output should therefore reach the user as structure rather than prose, and explanation depth should be progressive rather than fixed.

Three directions follow. First, objective comprehension measures and a deliberately incorrect-assessment condition, to determine whether the comprehension gain corresponds to calibrated

trust or to fluent acceptance. Second, a counterbalanced study with matched applications across conditions and a broader age range. Third, an interactive prototype supporting progressive disclosure, which would allow the efficiency measures—time to decision, depth of information sought—that a static stimulus could not provide.

REFERENCES

- [1] Hazim Almuhammedi, Florian Schaub, Norman Sadeh, Idris Adjerid, Alessandro Acquisti, Joshua Gluck, Lorrie Faith Cranor, and Yuvraj Agarwal. 2015. Your Location Has Been Shared 5,398 Times! A Field Study on Mobile App Privacy Nudging. In *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems (CHI '15)*. ACM, 787–796. <https://doi.org/10.1145/2702123.2702210>
- [2] Saleema Amershi, Dan Weld, Mihaela Vorvoreanu, Adam Fourney, Besmira Nushi, Penny Collisson, Jina Suh, Shamsi Iqbal, Paul N. Bennett, Kori Inkpen, Jaime Teevan, Ruth Kikin-Gil, and Eric Horvitz. 2019. Guidelines for Human-AI Interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems (CHI '19)*. ACM, 1–13. <https://doi.org/10.1145/3290605.3300233>
- [3] Daniel Arp, Michael Spreitzenbarth, Malte Hübner, Hugo Gascon, and Konrad Rieck. 2014. DREBIN: Effective and Explainable Detection of Android Malware in Your Pocket. In *Proceedings of the 21st Annual Network and Distributed System Security Symposium (NDSS '14)*. Internet Society. <https://doi.org/10.14722/ndss.2014.23247>
- [4] Gagan Bansal, Tongshuang Wu, Joyce Zhou, Raymond Fok, Besmira Nushi, Ece Kamar, Marco Tulio Ribeiro, and Daniel Weld. 2021. Does the Whole Exceed Its Parts? The Effect of AI Explanations on Complementary Team Performance. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI '21)*. ACM, 1–16. <https://doi.org/10.1145/3411764.3445717>
- [5] Virginia Braun and Victoria Clarke. 2006. Using Thematic Analysis in Psychology. *Qualitative Research in Psychology* 3, 2 (2006), 77–101. <https://doi.org/10.1191/1478088706qp0630a>
- [6] Alan Cooper, Robert Reimann, David Cronin, and Christopher Noessel. 2014. *About Face: The Essentials of Interaction Design* (4th ed.). John Wiley & Sons, Indianapolis, IN, USA.
- [7] Finale Doshi-Velez and Been Kim. 2017. Towards a Rigorous Science of Interpretable Machine Learning. *arXiv preprint arXiv:1702.08608* (2017).
- [8] Adrienne Porter Felt, Elizabeth Ha, Serge Egelman, Ariel Haney, Erika Chin, and David Wagner. 2012. Android Permissions: User Attention, Comprehension, and Behavior. In *Proceedings of the Eighth Symposium on Usable Privacy and Security (SOUPS '12)*. ACM, 1–14. <https://doi.org/10.1145/2335356.2335360>
- [9] Harmanpreet Kaur, Harsha Nori, Samuel Jenkins, Rich Caruana, Hanna Wallach, and Jennifer Wortman Vaughan. 2020. Interpreting Interpretability: Understanding Data Scientists' Use of Interpretability Tools for Machine Learning. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems (CHI '20)*. ACM, 1–14. <https://doi.org/10.1145/3313831.3376219>
- [10] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. 2017. LightGBM: A Highly Efficient Gradient Boosting Decision Tree. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*. 3146–3154.
- [11] Patrick Gage Kelley, Joanna Bresee, Lorrie Faith Cranor, and Robert W. Reeder. 2009. A “Nutrition Label” for Privacy. In *Proceedings of the 5th Symposium on Usable Privacy and Security (SOUPS '09)*. ACM, 1–12. <https://doi.org/10.1145/1572532.1572538>
- [12] Patrick Gage Kelley, Lorrie Faith Cranor, and Norman Sadeh. 2013. Privacy as Part of the App Decision-Making Process. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '13)*. ACM, 3393–3402. <https://doi.org/10.1145/2470654.2466466>
- [13] Q. Vera Liao, Daniel Gruen, and Sarah Miller. 2020. Questioning the AI: Informing Design Practices for Explainable AI User Experiences. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems (CHI '20)*. ACM, 1–15. <https://doi.org/10.1145/3313831.3376590>
- [14] Jiali Lin, Shahriyar Amini, Jason I. Hong, Norman Sadeh, Janne Lindqvist, and Joy Zhang. 2012. Expectation and Purpose: Understanding Users' Mental Models of Mobile App Privacy through Crowdsourcing. In *Proceedings of the 2012 ACM Conference on Ubiquitous Computing (UbiComp '12)*. ACM, 501–510. <https://doi.org/10.1145/2370216.2370290>
- [15] Zachary C. Lipton. 2018. The Myths of Model Interpretability. *Queue* 16, 3 (2018), 31–57. <https://doi.org/10.1145/3236386.3241340>
- [16] Bin Liu, Mads Schaarup Andersen, Florian Schaub, Hazim Almuhammedi, Shikun Zhang, Norman Sadeh, Yuvraj Agarwal, and Alessandro Acquisti. 2016. Follow My Recommendations: A Personalized Privacy Assistant for Mobile App Permissions. In *Proceedings of the Twelfth Symposium on Usable Privacy and Security (SOUPS '16)*. USENIX Association, 27–41.
- [17] Scott M. Lundberg, Gabriel Erion, Hugh Chen, Alex DeGrave, Jordan M. Prutkin, Bala Nair, Ronit Katz, Jonathan Himmelfarb, Nisha Bansal, and Su-In Lee. 2020. From Local Explanations to Global Understanding with Explainable AI for Trees. *Nature Machine Intelligence* 2, 1 (2020), 56–67. <https://doi.org/10.1038/s42256-019-0138-9>
- [18] Scott M. Lundberg and Su-In Lee. 2017. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*. 4765–4774.
- [19] Tim Miller. 2019. Explanation in Artificial Intelligence: Insights from the Social Sciences. *Artificial Intelligence* 267 (2019), 1–38. <https://doi.org/10.1016/j.artint.2018.07.007>
- [20] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. 2011. Machine Learning in Python. *Journal of Machine Learning Research* 12 (2011), 2825–2830.
- [21] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. “Why Should I Trust You?”: Explaining the Predictions of Any Classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '16)*. ACM, 1135–1144. <https://doi.org/10.1145/2939672.2939778>
- [22] Helen Sharp, Jennifer Preece, and Yvonne Rogers. 2019. *Interaction Design: Beyond Human-Computer Interaction* (5th ed.). John Wiley & Sons, Indianapolis, IN, USA.
- [23] Ben Shneiderman. 2020. Bridging the Gap Between Ethics and Practice: Guidelines for Reliable, Safe, and Trustworthy Human-Centered AI Systems. *ACM Transactions on Interactive Intelligent Systems* 10, 4 (2020), 1–31. <https://doi.org/10.1145/3419764>
- [24] Ben Shneiderman. 2022. *Human-Centered AI*. Oxford University Press, Oxford, UK.
- [25] Wei Wang, Xing Wang, Dawei Feng, Jiqiang Liu, Zhen Han, and Xiangliang Zhang. 2014. Exploring Permission-Induced Risk in Android Applications for Malicious Application Detection. *IEEE Transactions on Information Forensics and Security* 9, 11 (2014), 1869–1882. <https://doi.org/10.1109/TIFS.2014.2353996>
- [26] Frank Wilcoxon. 1945. Individual Comparisons by Ranking Methods. *Biometrics Bulletin* 1, 6 (1945), 80–83. <https://doi.org/10.2307/3001968>
- [27] Suleiman Y. Yerima and Sakir Sezer. 2019. DroidFusion: A Novel Multilevel Classifier Fusion Approach for Android Malware Detection. *IEEE Transactions on Cybernetics* 49, 2 (2019), 453–466. <https://doi.org/10.1109/TCYB.2017.2777960>
- [28] Shikun Zhang and Norman Sadeh. 2023. Do Privacy Labels Answer Users' Privacy Questions?. In *Workshop on Usable Security and Privacy (USEC '23)*. Internet Society.